




Generative AI LangChain

ASHIMA MALIK

 ashimamalik58@gmail.com



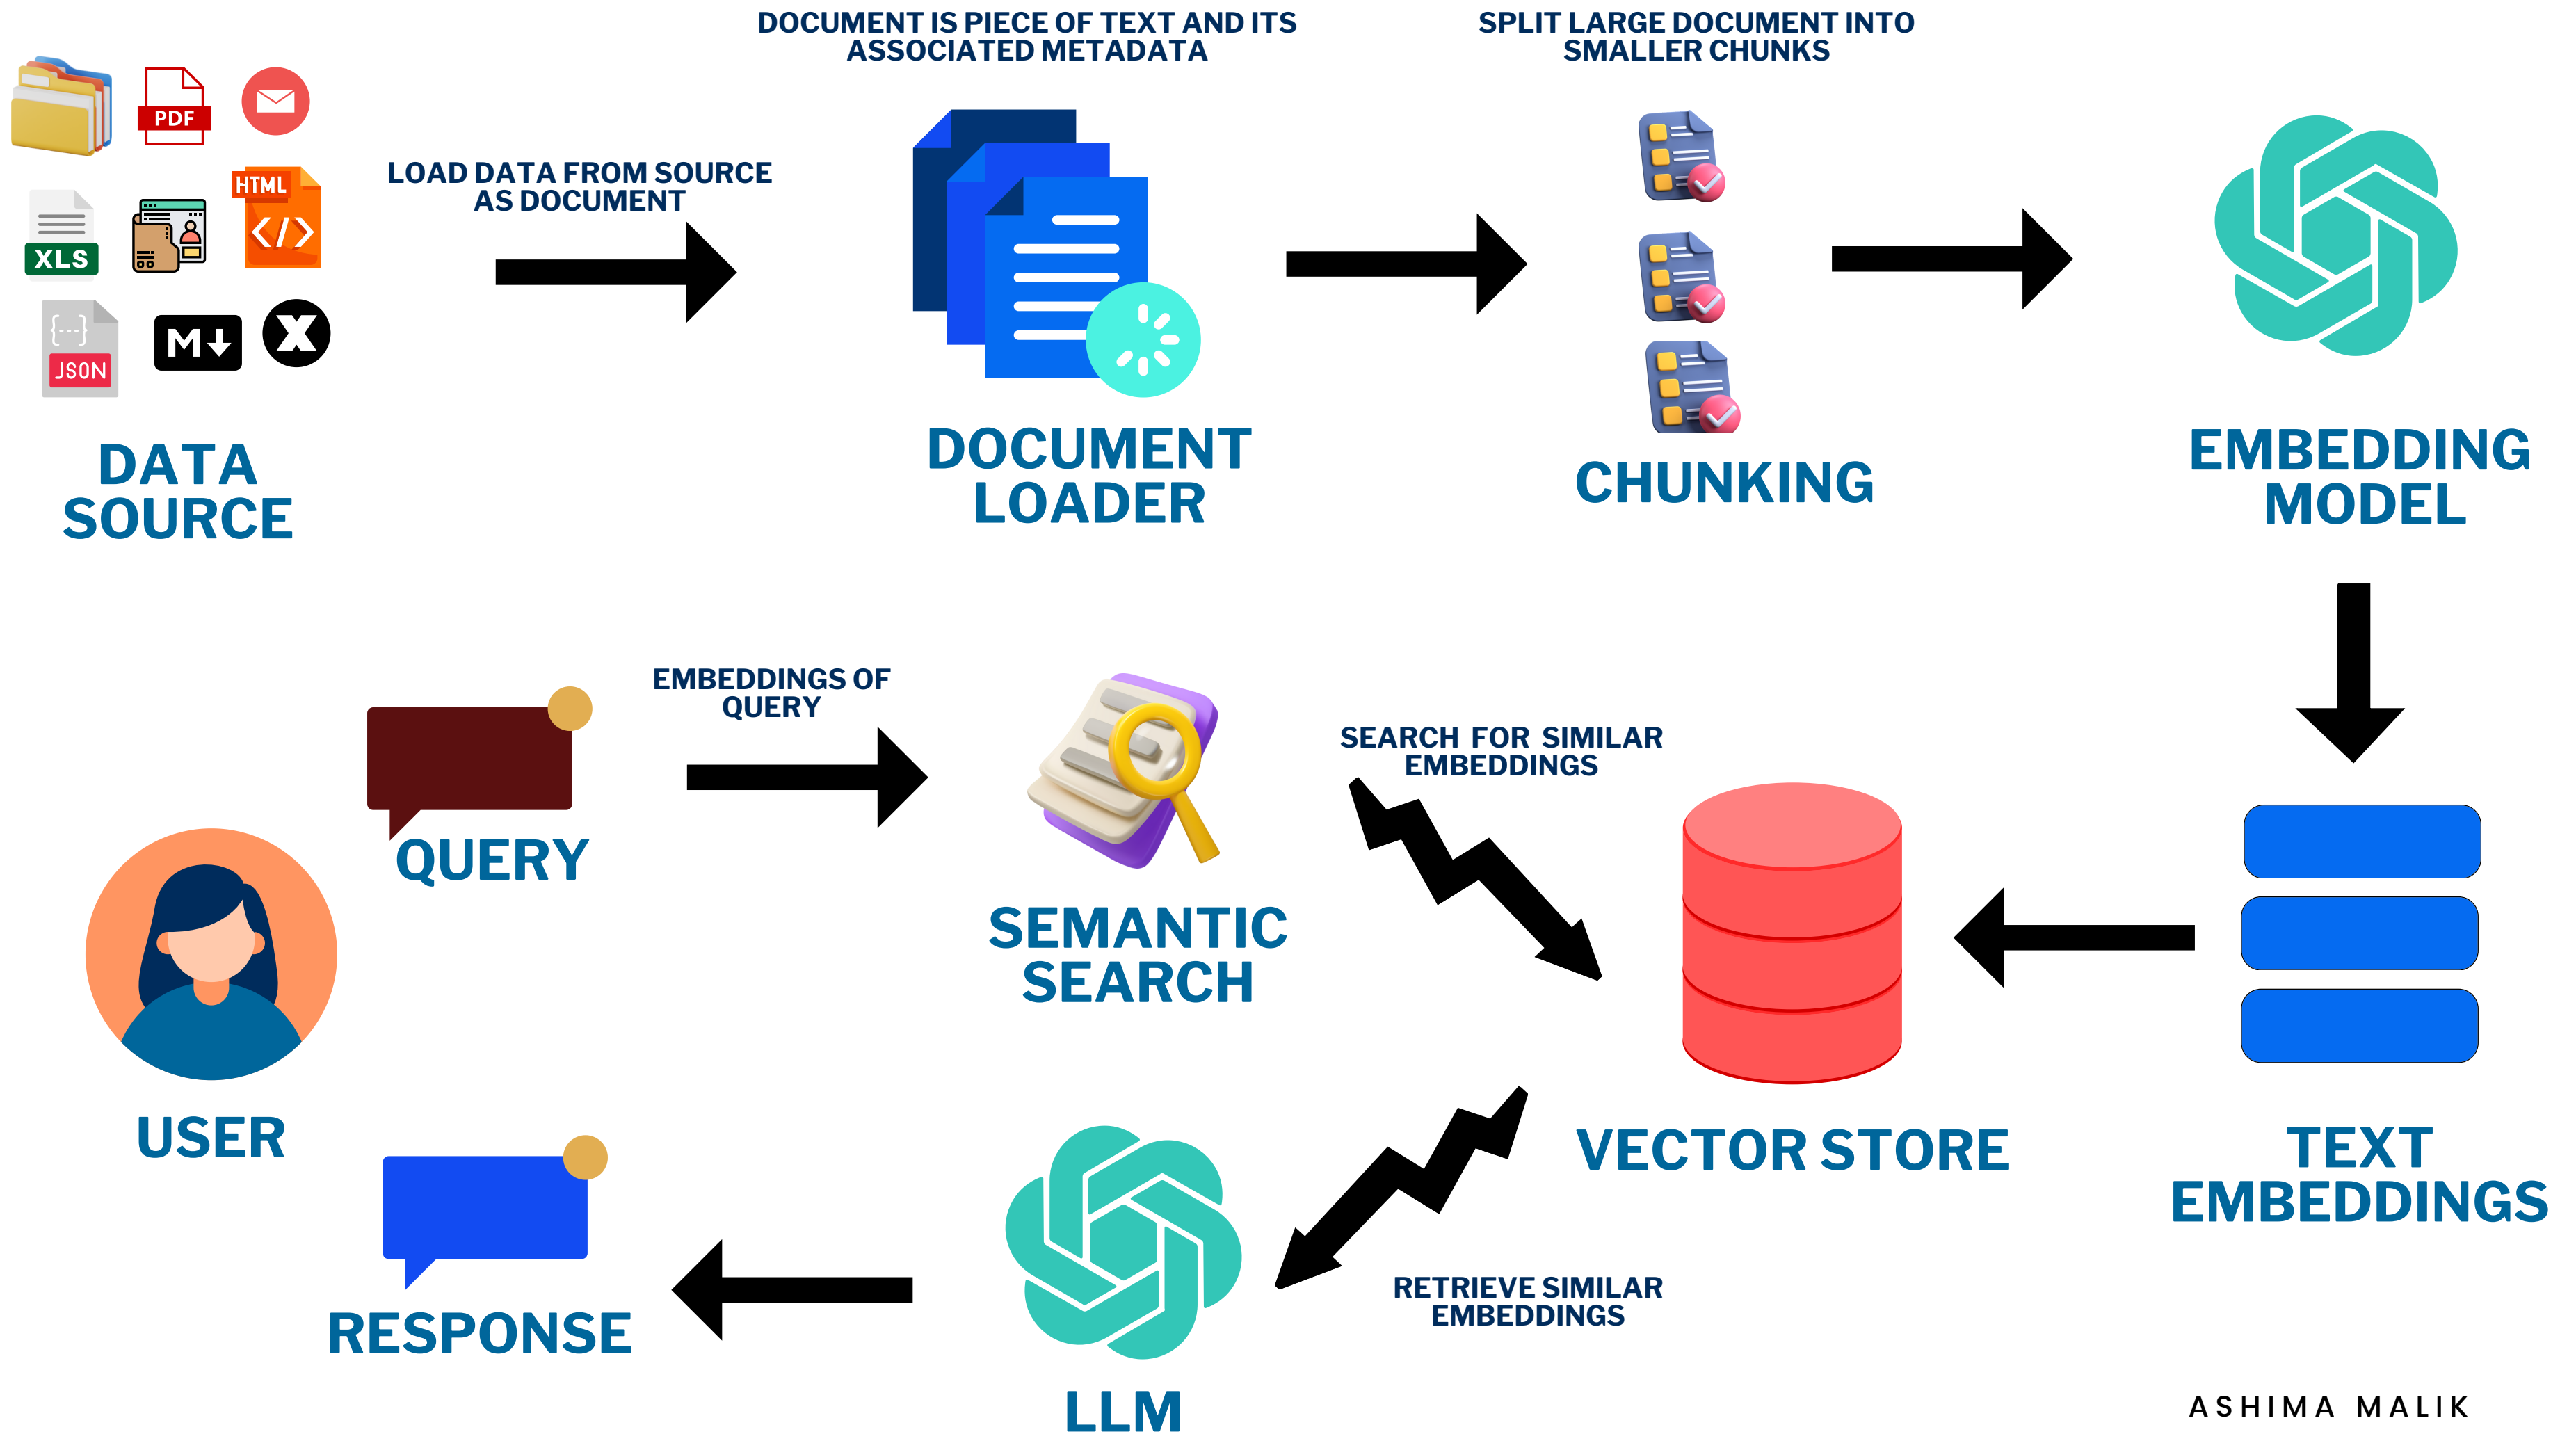
Retrieval



Limited by training data, many LLMs struggle with user-specific information. Retrieval-Augmented Generation (RAG) tackles this challenge by retrieving relevant external data and incorporating it into the generation process alongside the LLM's internal knowledge.

LLMs benefit from information beyond their training data.

RAG helps bridge this gap by retrieving relevant external data for the LLM to utilize during generation.



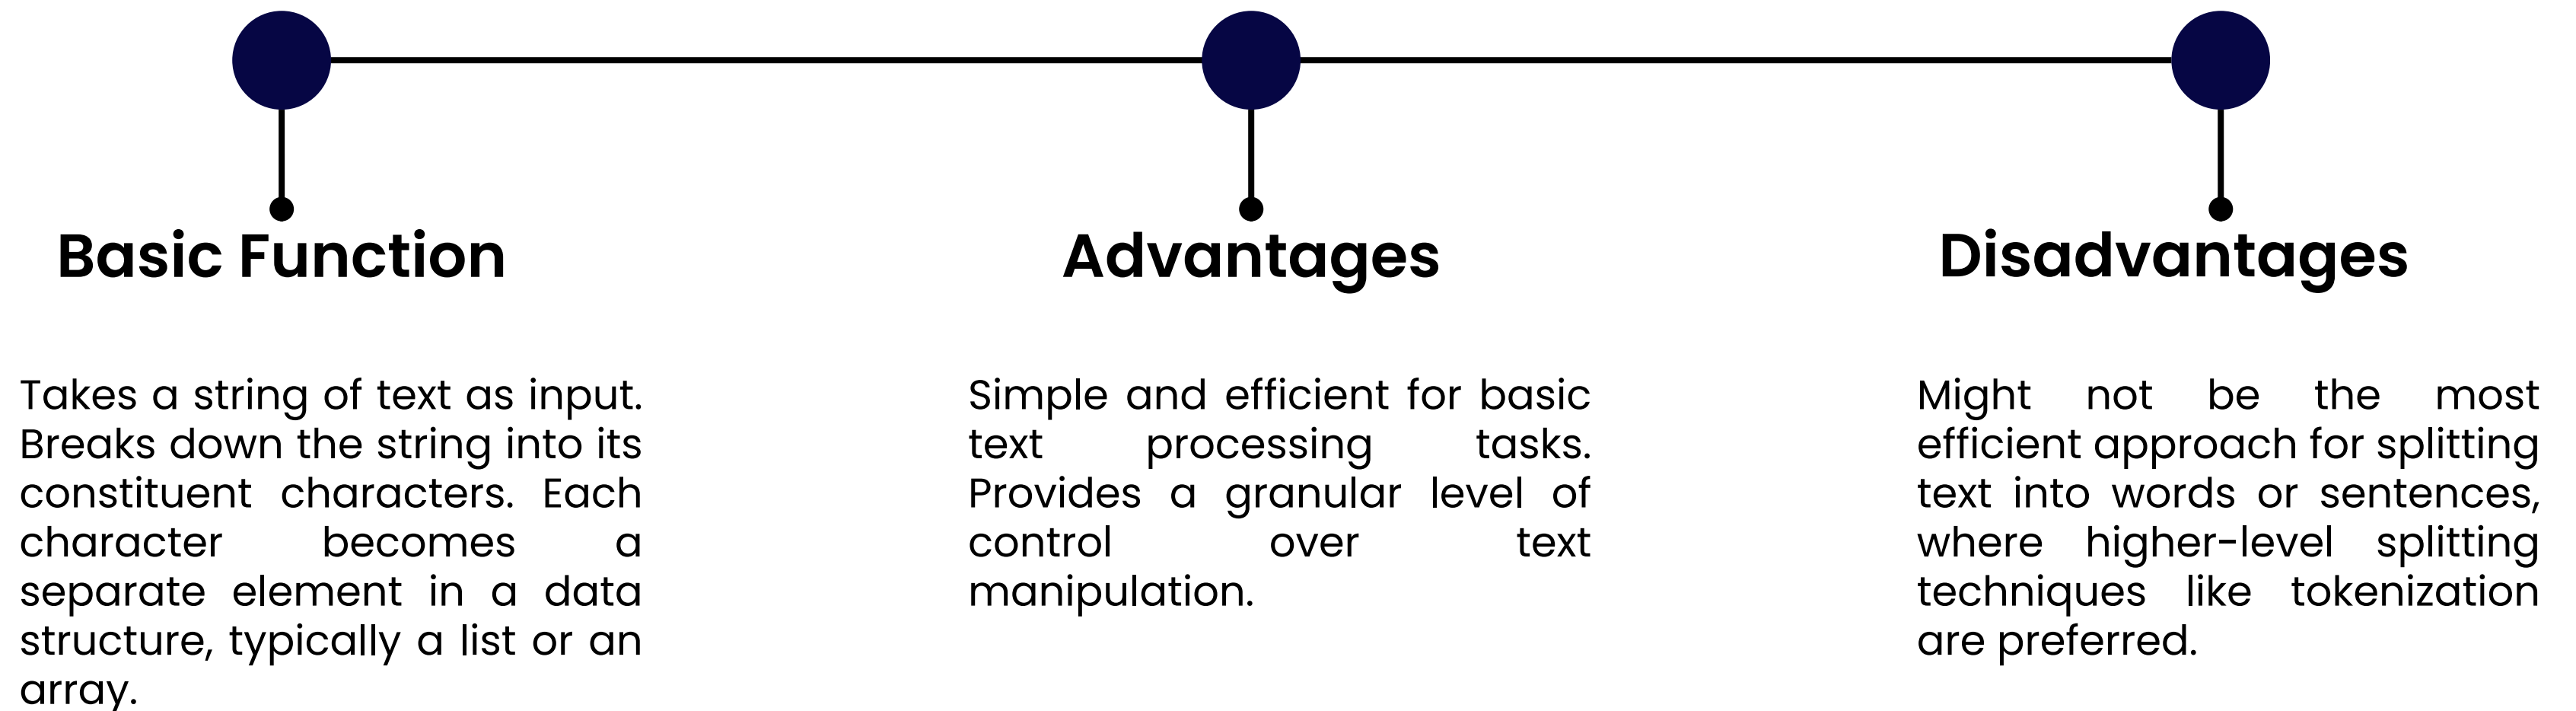
DOCUMENT LOADERS

| Document Loader (Import Statement) | Function |
|--|---|
| TextLoader (from langchain_community.document_loaders import TextLoader) | Loads text content from a file |
| CSVLoader (from langchain_community.document_loaders.csv_loader import CSVLoader) | Loads data from a CSV file into Document objects, with each row becoming a Document |
| JSONLoader (from langchain_community.document_loaders import JSONLoader) | Reads JSON data from a file where each line represents a JSON object, converting each object into a Document. |
| HTMLLoader (from langchain_community.document_loaders import UnstructuredHTMLLoader) | Fetches the text content and basic metadata from a website URL, creating a Document. |
| MarkdownLoader (from langchain_community.document_loaders import UnstructuredMarkdownLoader) | Loads the text content from a Markdown file, converting it into a Document. |
| PDFLoader (from langchain.document_loaders import PyPDFLoader) | Extracts text content from a PDF file, creating a Document (may require additional libraries). |
| DirectoryLoader (from langchain.document_loaders import Directory_Loader) | Loads all files within a specified directory as Documents, using appropriate loaders based on file extensions (.txt, .csv, etc.). |

TEXT SPLITTERS

Character Text Splitter

Character Splitting, in the context of text processing, refers to the process of dividing a string of text into individual characters. It's a fundamental operation used in various text manipulation tasks, data analysis, and machine learning applications.



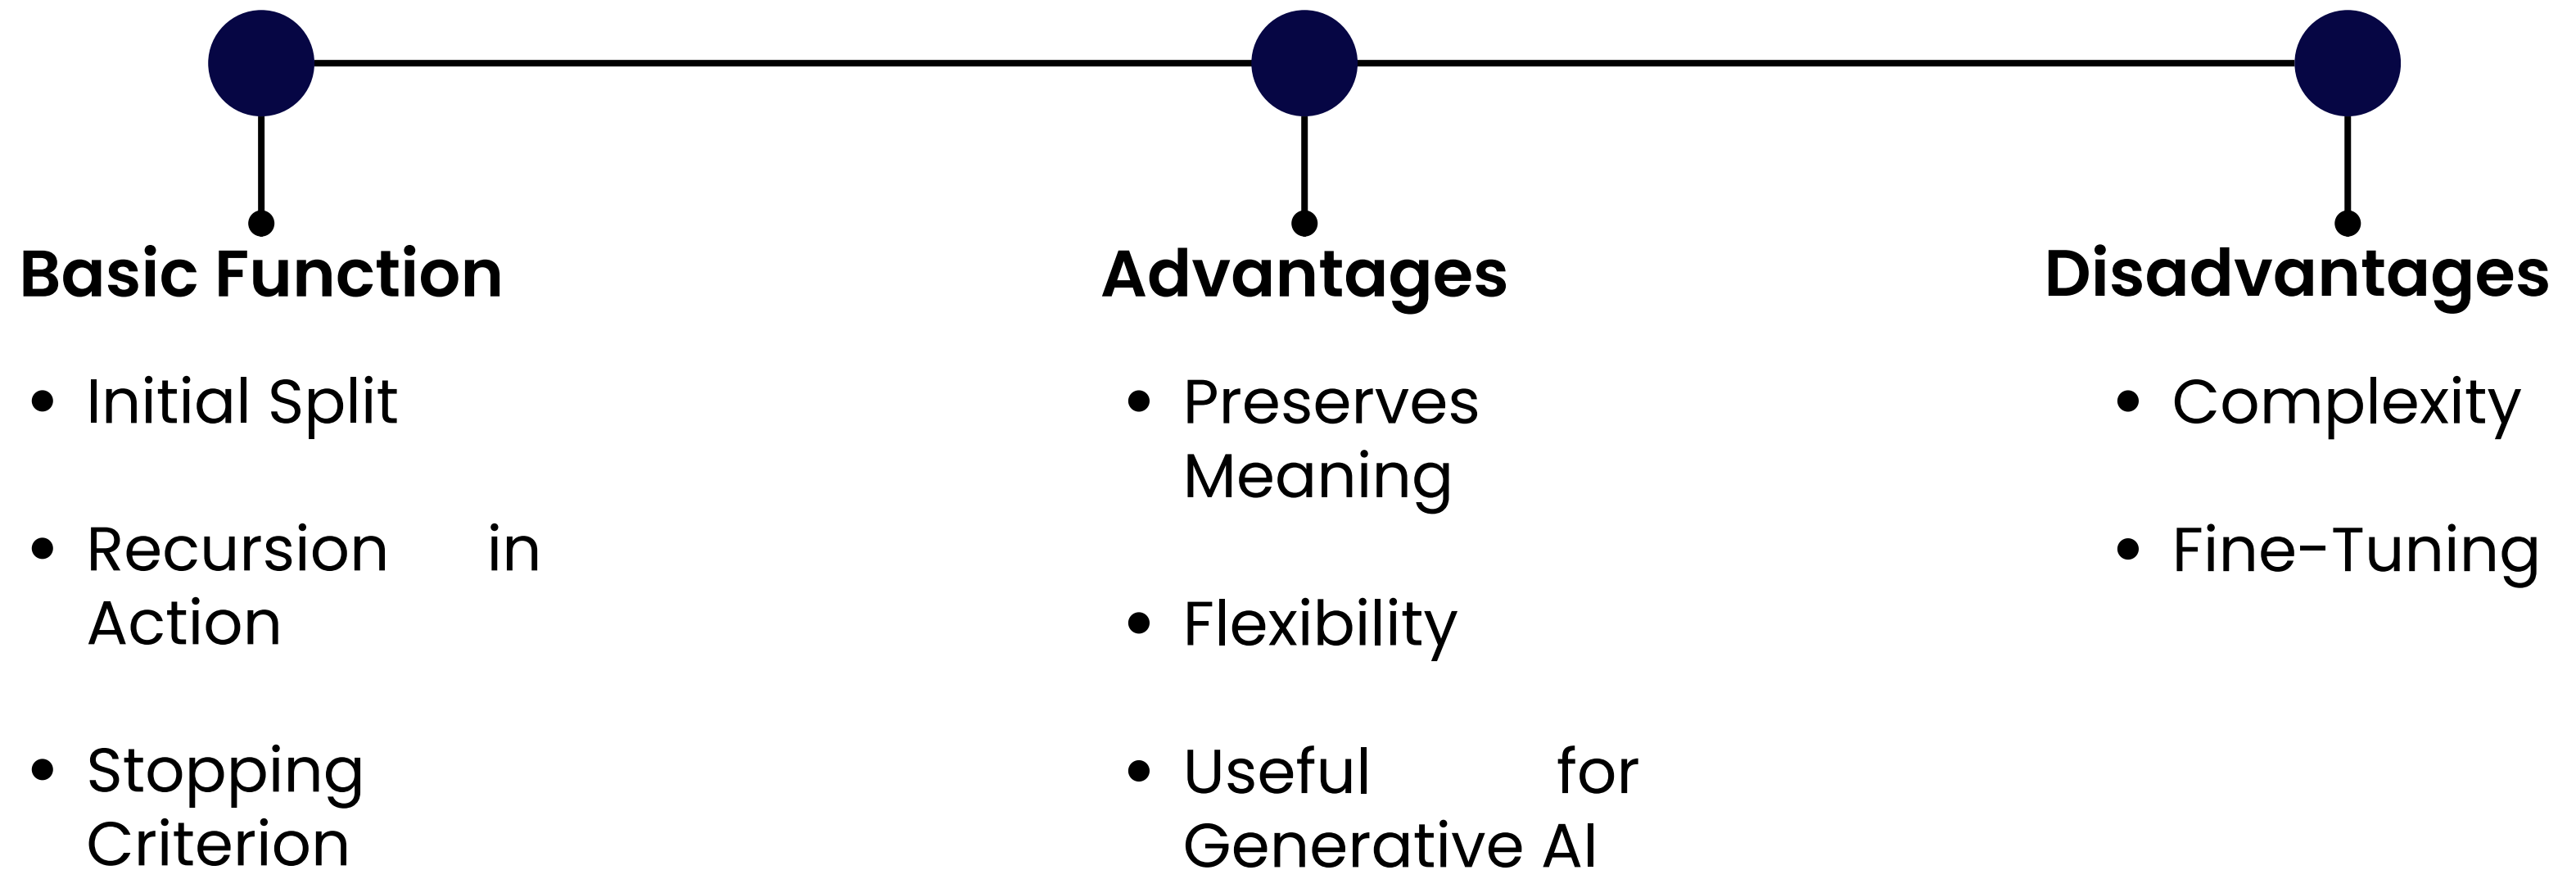
Recursive Character Text Splitting

"\n\n" - Double new line, or most commonly paragraph breaks

"\n" - New lines

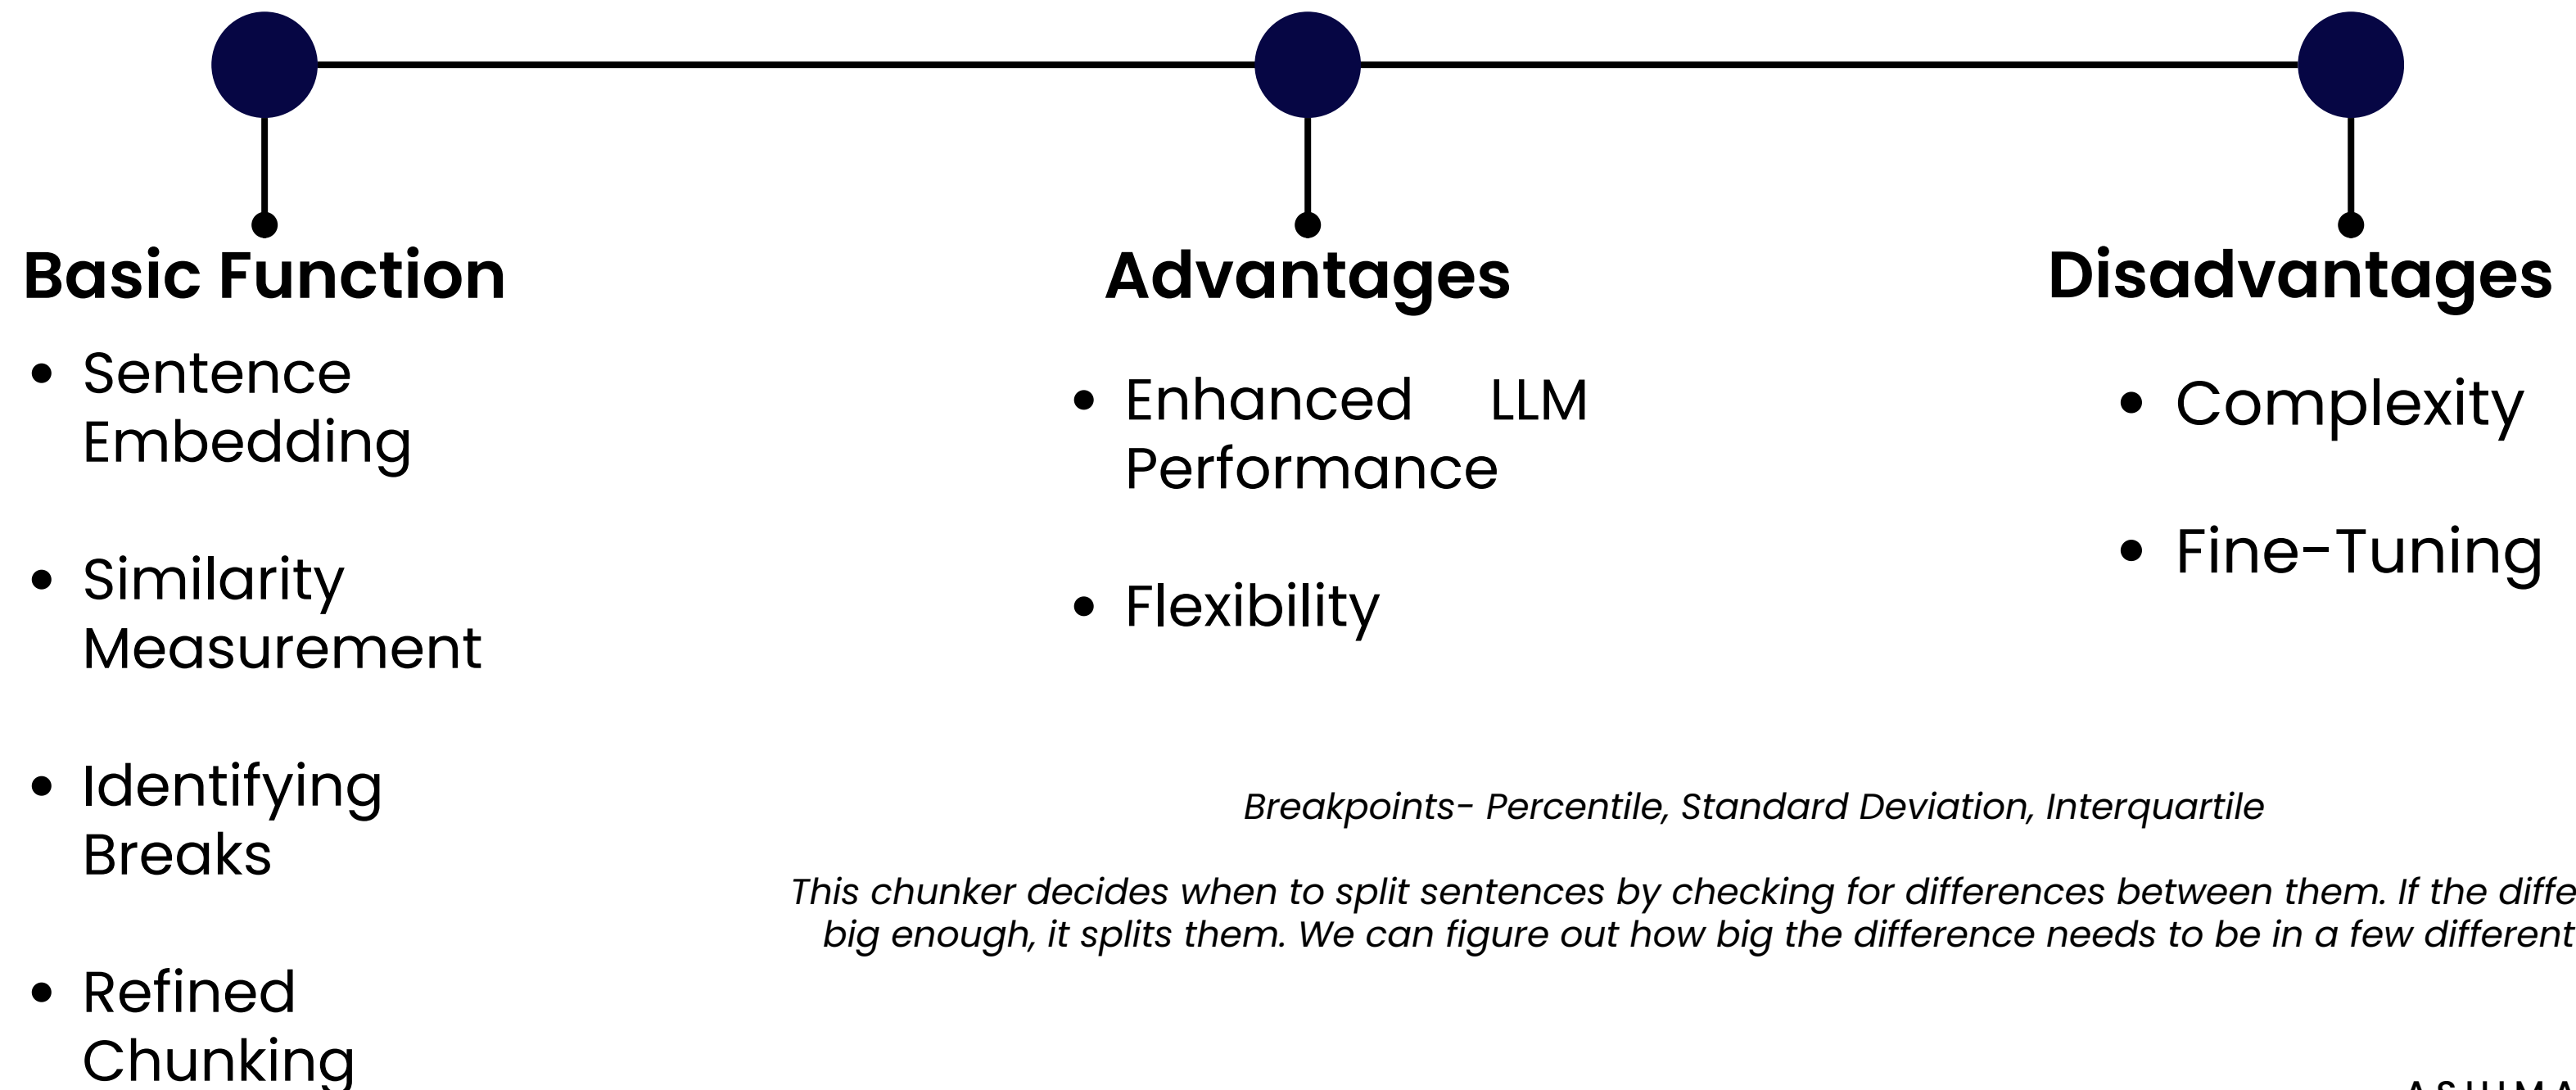
" " - Spaces

"" - Characters



Semantic Chunking

LangChain utilizes techniques like sentence embedding to determine semantic similarity between different parts of the text. Sentence embedding involves converting sentences into numerical vectors that capture their meaning.



TEXT EMBEDDING MODELS

TEXT

Actions speaks.

Quick Learner.

He is perfect.

MODELS



EMBEDDINGS

[0.2, 0.4, -0.1.....]

[0.1, 0.3, 0.2]

[-0.1, 0.4, -0.3]

VECTOR STORES & RETRIEVERS

WORD EMBEDDINGS

****Step1: Collect data****

It was the best of times,
it was the worst of times,
it was the age of wisdom,
it was the age of foolishness!

****Step2: Design the vocabulary****

Now we can make a list of all of the words in our model vocabulary.
The unique words here (ignoring case and punctuation) are:

"it"
"was"
"the"
"best"
"of"
"times"
"worst"
"age"
"wisdom"
"foolishness"

This is a vocabulary of 10 words from a corpus containing 24 words.

****Step3: Create document vectors****


it was the best of time worst age wisdom foolishness

| | | | | | | | | | | |
|-----------------------------------|---|---|---|---|---|---|---|---|---|---|
| "it was the worst of times" = | 1 | 1 | 1 | 0 | 1 | 1 | 1 | 0 | 0 | 0 |
| "it was the age of wisdom" = | 1 | 1 | 1 | 0 | 1 | 0 | 0 | 1 | 1 | 0 |
| "it was the age of foolishness" = | 1 | 1 | 1 | 0 | 1 | 0 | 0 | 1 | 0 | 1 |

WORD EMBEDDINGS

king

[0.50451, 0.68607, -0.59517, -0.022801, 0.60046, -0.13498, -0.08813, 0.47377, -0.61798, -0.31012, -0.076666, 1.493, -0.034189, -0.98173, 0.68229, 0.81722, -0.51874, -0.31503, -0.55809, 0.66421, 0.1961, -0.13495, -0.11476, -0.30344, 0.41177, -2.223, -1.0756, -1.0783, -0.34354, 0.33505, 1.9927, -0.04234, -0.64319, 0.71125, 0.49159, 0.16754, 0.34344, -0.25663, -0.8523, 0.1661, 0.40102, 1.1685, -1.0137, -0.21585, -0.15155, 0.78321, -0.91241, -1.6106, -0.64426, -0.51042]



“king”



“Man”



“Woman”



Image Source: <http://jalammar.github.io/illustrated-word2vec/>

WORD EMBEDDINGS



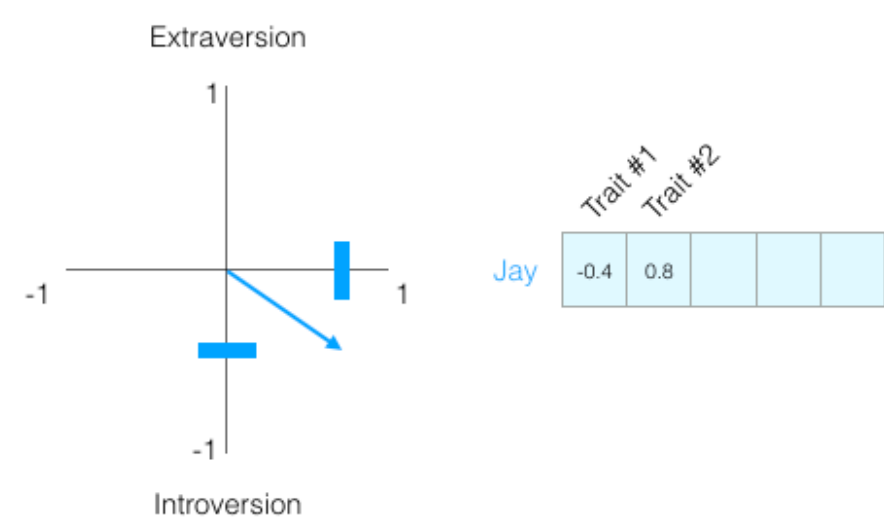



Image Source: <http://jalammar.github.io/illustrated-word2vec/>

WORD EMBEDDINGS



$\text{cosine_similarity}(\text{Jay}, \text{Person \#1}) = 0.87$ ✓

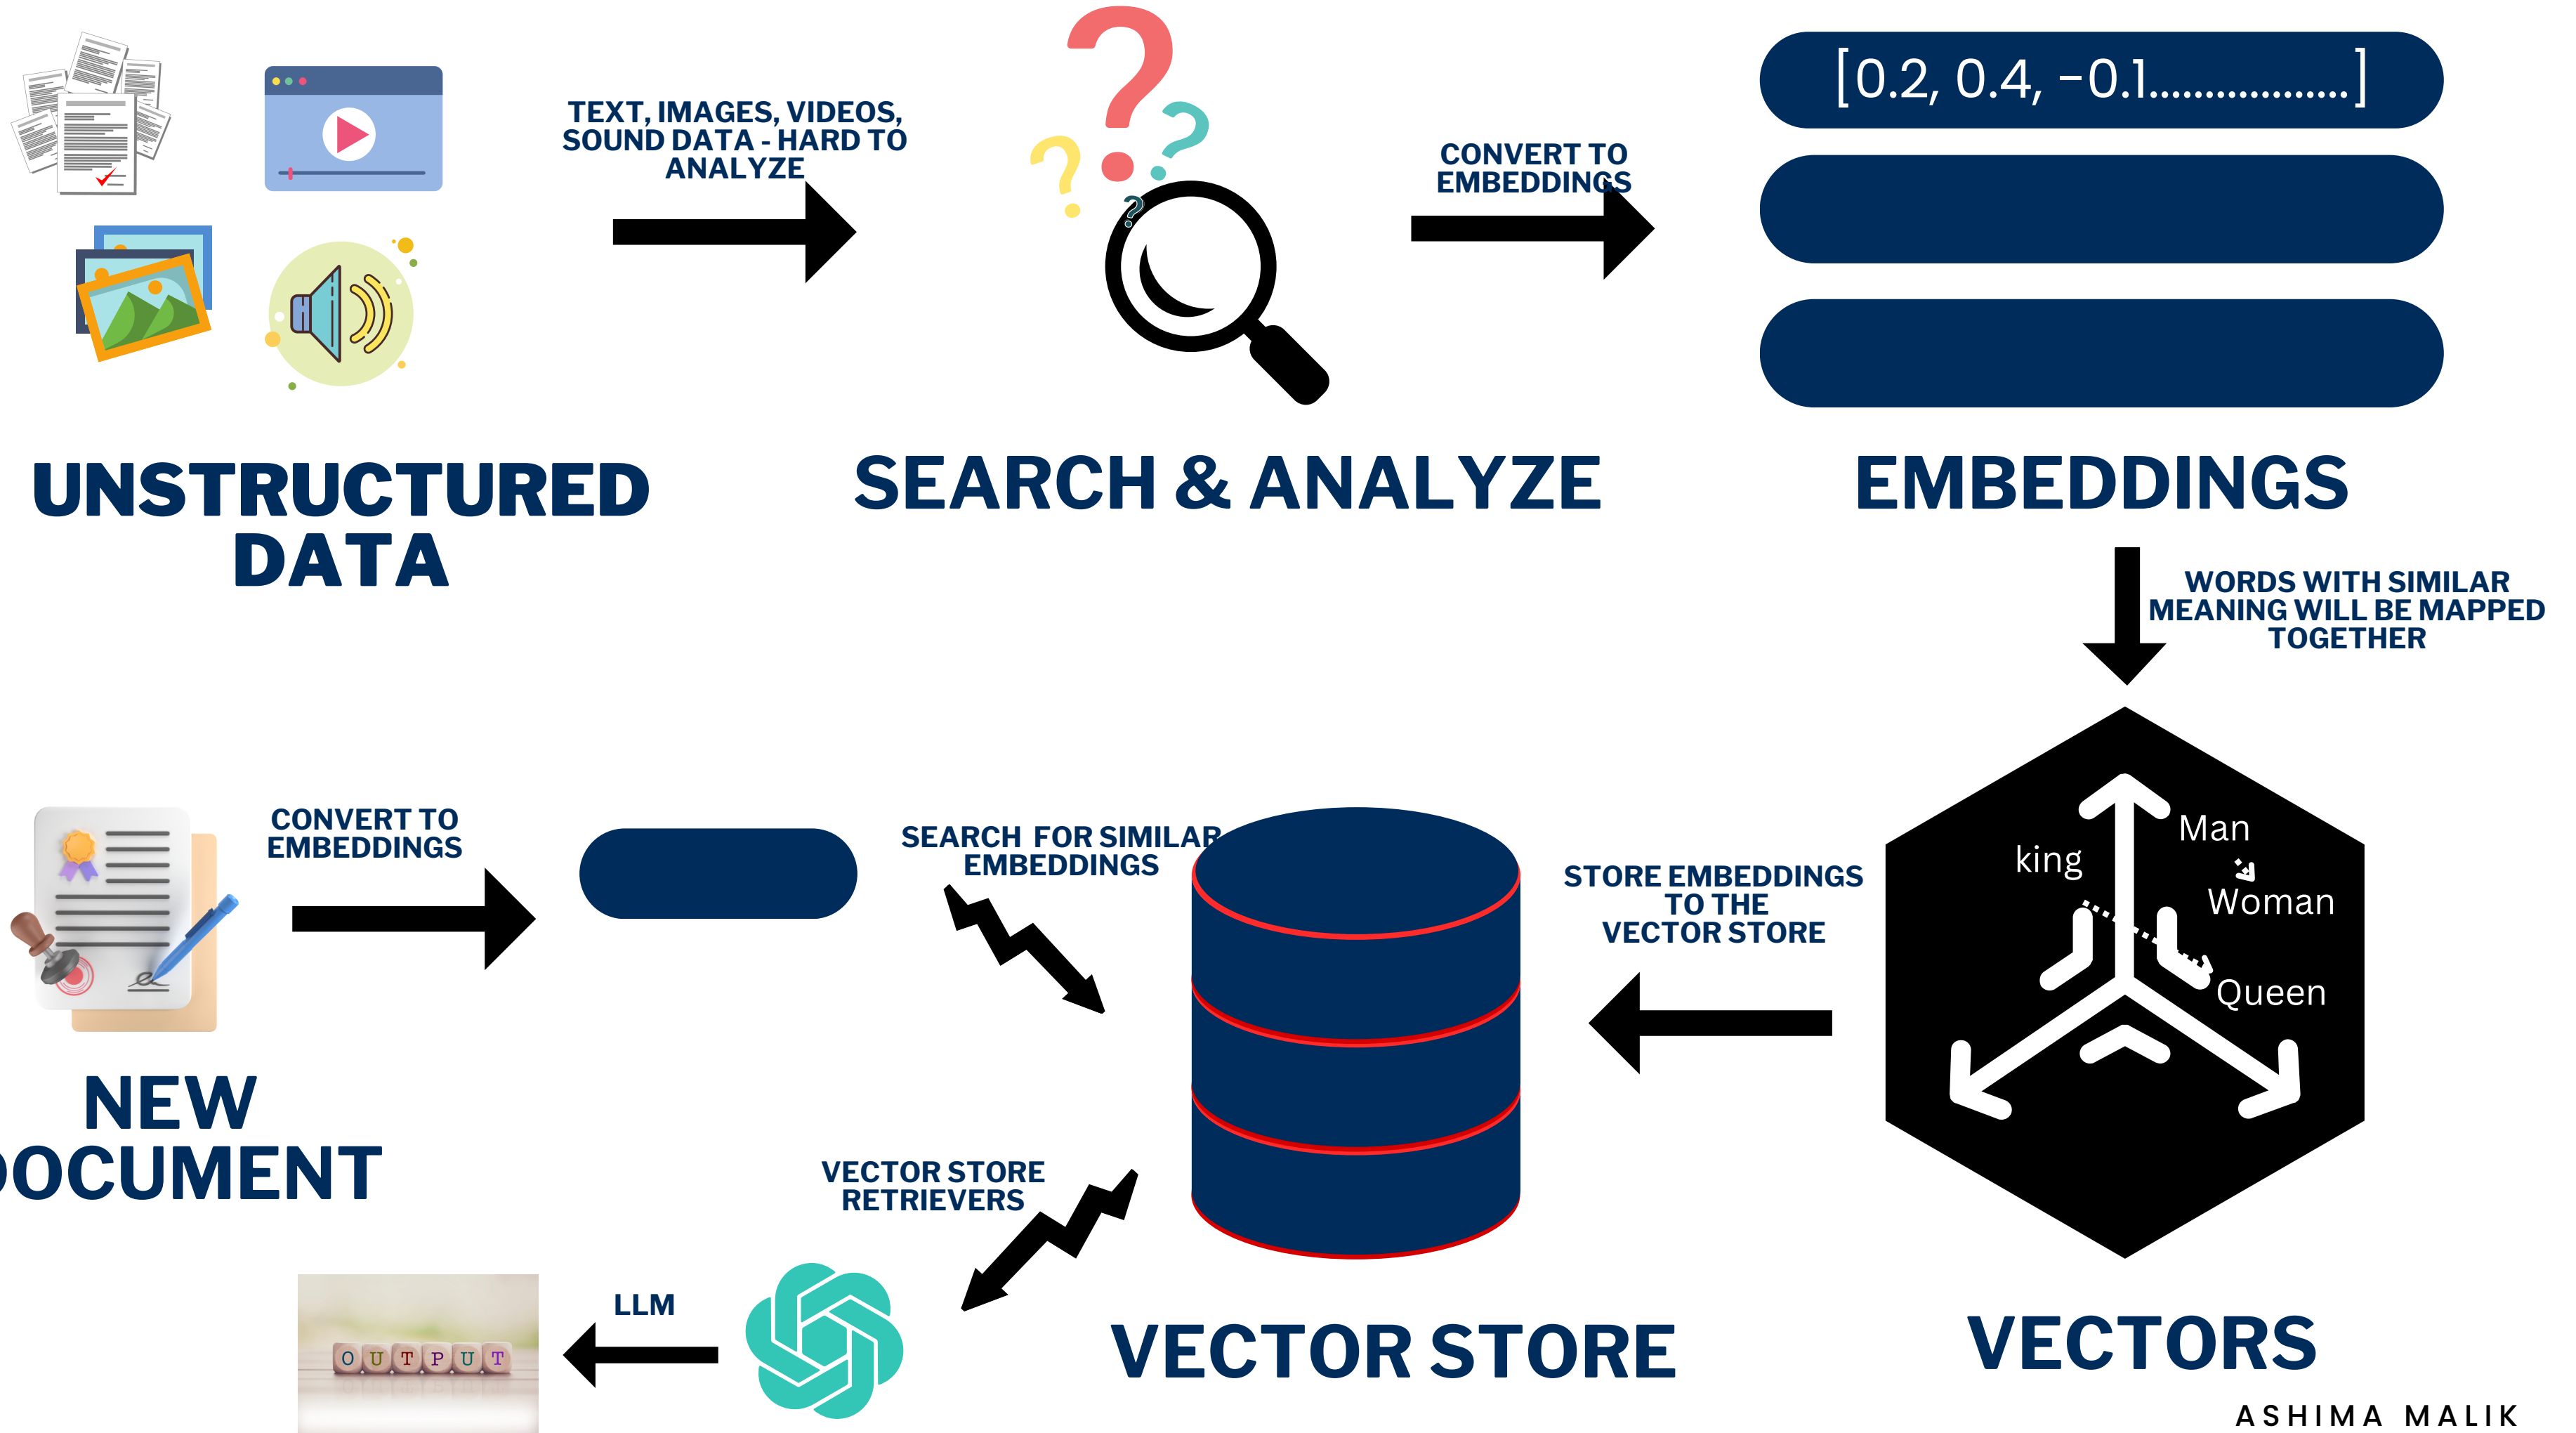
$\text{cosine_similarity}(\text{Jay}, \text{Person \#2}) = -0.20$



$\text{cosine_similarity}(\text{Jay}, \text{Person \#1}) = 0.66$ ✓

$\text{cosine_similarity}(\text{Jay}, \text{Person \#2}) = -0.37$


Image Source: <http://jalammar.github.io/illustrated-word2vec/>



| Name | Index Type | Uses an LLM | When to Use | Description |
|---------------------------|------------------------------|-------------|---|--|
| Vectorstore | Vectorstore | No | If you are just getting started and looking for something quick and easy. | This is the simplest method and the one that is easiest to get started with. It involves creating embeddings for each piece of text. |
| ParentDocument | Vectorstore + Document Store | No | If your pages have lots of smaller pieces of distinct information that are best indexed by themselves, but best retrieved all together. | This involves indexing multiple chunks for each document. Then you find the chunks that are most similar in embedding space, but you retrieve the whole parent document and return that (rather than individual chunks). |
| Multi Vector | Vectorstore + Document Store | Sometimes | If you are able to extract information from documents that you think is more relevant to index than the text itself. | This involves creating multiple vectors for each document. Each vector could be created in a myriad of ways - examples include summaries of the text and hypothetical questions. |
| Self Query | Vectorstore | Yes | If users are asking questions that are better answered by fetching documents based on metadata rather than similarity with the text. | This uses an LLM to transform user input into two things: (1) a string to look up semantically, (2) a metadata filter to go along with it. This is useful because oftentimes questions are about the METADATA of documents (not the content itself). |
| Contextual Compression | Any | Sometimes | If you are finding that your retrieved documents contain too much irrelevant information and are distracting the LLM. | This puts a post-processing step on top of another retriever and extracts only the most relevant information from retrieved documents. This can be done with embeddings or an LLM. |
| Time-Weighted Vectorstore | Vectorstore | No | If you have timestamps associated with your documents, and you want to retrieve the most recent ones | This fetches documents based on a combination of semantic similarity (as in normal vector retrieval) and recency (looking at timestamps of indexed documents) |
| Multi-Query Retriever | Any | Yes | If users are asking questions that are complex and require multiple pieces of distinct information to respond | This uses an LLM to generate multiple queries from the original one. This is useful when the original query needs pieces of information about multiple topics to be properly answered. By generating multiple queries, we can then fetch documents for each of them. |
| Ensemble | Any | No | If you have multiple retrieval methods and want to try combining them. | This fetches documents from multiple retrievers and then combines them. |
| Long-Context Reorder | Any | No | If you are working with a long-context model and noticing that it's not paying attention to information in the middle of retrieved documents. | This fetches documents from an underlying retriever, and then reorders them so that the most similar are near the beginning and end. |

Indexing

The indexing API enables you to import and maintain documents from various sources within a vector store efficiently.



RecordManager keeps track of document writes into the vector store.

Its primary functions include preventing duplicate content from being written to the vector store, avoiding rewriting content that remains unchanged, and preventing the recomputation of embeddings for content that has not been altered.